AgentRQ WHAT IT SOLVES TODAY

Hand AI agents real work.
Stay in control.

Self-learning task manager and orchestration layer for AI agents: any harness, any machine, one board.

01 DISTRIBUTED SKILLS

- Playbooks copied to every box

- Stored once, served over MCP to every agent, machine and harness. Edit once, and every workspace reads the new copy.

02 ONE TASK BOARD

- Babysitting every terminal

- Agents pick up tasks, report status and ask you questions on one live board, with Allow and Deny right on the task.

03 APPROVALS + YOLO

- Safe or fast, never both

- Approve per task, or turn on YOLO for a task or workspace and still keep a record of every call it made.

04 TRAJECTORY

- No idea what an agent ran

- Every tool call and message, searchable, plus a timeline of time worked, blocked and waiting on you.

05 ISOLATED WORKSPACES

- Agents step on each other

- Each workspace has its own MCP address, memory and folder. Forks run in parallel in their own git worktree.

06 CRON + HOOKS

- Work stops when you look away

- Cron runs, events and drag and drop workflows. A failed CI run on GitHub arrives as a task for an agent.

07 MEMORY + PANO

- Agents forget everything

- Workspace memory outlives a task, and shared Pano docs pass plans and checklists between agents.

08 LIVE TERMINALS

- Agents tied to one laptop

- Agents run on your machines, in live terminals you open from where you go: web, desktop and mobile.

09 ANY AGENT

- Every harness is different

- One MCP surface for any MCP or ACP agent: Claude Code, Codex, Antigravity, Hermes, OpenClaw and more.

10 SITE TOOLS

- Logged-in sites off limits

- The Chrome extension lends a site's WebMCP tools to agents, run in your own signed-in browser, with your approval.

CLAUDE CODEX AGY HERMES OPENCLAW

hi@agentrq.com

Mustafa Turan · Founder @ Contextual, Inc.